

US House Price Prediction

Table of Contents

This project contains:

- About the data
- Exploratory Data Analysis
- Preparing the dataset for training
- Training and Validation set
- Random Forest Classifier
- Making Predictions
- Conclusion

About the data:

This dataset is from the Kaggle platform.

<https://www.kaggle.com/competitions/house-prices-advanced-regression-techniques>



The main motive of this project is to predict the price of a house information like its location, area, no. of rooms etc. Random Forest Classifier model will be implemented to predict the final price of each home.

```
# installing opendatasets library
!pip install opendatasets
```

Requirement already satisfied: opendatasets in /opt/conda/lib/python3.9/site-packages (0.1.22)

Requirement already satisfied: click in /opt/conda/lib/python3.9/site-packages (from opendatasets) (8.0.3)

Requirement already satisfied: tqdm in /opt/conda/lib/python3.9/site-packages (from opendatasets) (4.62.3)

Requirement already satisfied: kaggle in /opt/conda/lib/python3.9/site-packages (from opendatasets) (1.5.12)

Requirement already satisfied: requests in /opt/conda/lib/python3.9/site-packages (from kaggle->opendatasets) (2.26.0)

Requirement already satisfied: certifi in /opt/conda/lib/python3.9/site-packages (from kaggle->opendatasets) (2021.10.8)

Requirement already satisfied: python-dateutil in /opt/conda/lib/python3.9/site-packages (from kaggle->opendatasets) (2.8.2)

Requirement already satisfied: six>=1.10 in /opt/conda/lib/python3.9/site-packages (from kaggle->opendatasets) (1.16.0)

Requirement already satisfied: urllib3 in /opt/conda/lib/python3.9/site-packages (from

kaggle->opendatasets) (1.26.7)

Requirement already satisfied: python-slugify in /opt/conda/lib/python3.9/site-packages (from kaggle->opendatasets) (7.0.0)

Requirement already satisfied: text-unidecode>=1.3 in /opt/conda/lib/python3.9/site-packages (from python-slugify->kaggle->opendatasets) (1.3)

Requirement already satisfied: idna<4,>=2.5 in /opt/conda/lib/python3.9/site-packages (from requests->kaggle->opendatasets) (3.1)

Requirement already satisfied: charset-normalizer~=2.0.0 in /opt/conda/lib/python3.9/site-packages (from requests->kaggle->opendatasets) (2.0.0)

```
!pip install jovian
```

Requirement already satisfied: jovian in /opt/conda/lib/python3.9/site-packages (0.2.45)

Requirement already satisfied: uuid in /opt/conda/lib/python3.9/site-packages (from jovian) (1.30)

Requirement already satisfied: click in /opt/conda/lib/python3.9/site-packages (from jovian) (8.0.3)

Requirement already satisfied: requests in /opt/conda/lib/python3.9/site-packages (from jovian) (2.26.0)

Requirement already satisfied: pyyaml in /opt/conda/lib/python3.9/site-packages (from jovian) (6.0)

Requirement already satisfied: certifi>=2017.4.17 in /opt/conda/lib/python3.9/site-packages (from requests->jovian) (2021.10.8)

Requirement already satisfied: idna<4,>=2.5 in /opt/conda/lib/python3.9/site-packages (from requests->jovian) (3.1)

Requirement already satisfied: urllib3<1.27,>=1.21.1 in /opt/conda/lib/python3.9/site-packages (from requests->jovian) (1.26.7)

Requirement already satisfied: charset-normalizer~=2.0.0 in /opt/conda/lib/python3.9/site-packages (from requests->jovian) (2.0.0)

```
# importing necessary libraries
import opendatasets as od
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib
import pandas as pd
import numpy as np
import jovian
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.ensemble import RandomForestClassifier
%matplotlib inline
matplotlib.rcParams['figure.facecolor']='white'
```

```
data_url='https://www.kaggle.com/competitions/house-prices-advanced-regression-techniques'
```

```
od.download(data_url)
```

Skipping, found downloaded files in "./house-prices-advanced-regression-techniques"
(use force=True to force download)

```
data_dir='./house-prices-advanced-regression-techniques'
```

```
import os
```

```
os.listdir(data_dir)
```

```
['test.csv', 'train.csv', 'sample_submission.csv', 'data_description.txt']
```

```
test_csv=data_dir+'/test.csv'  
train_csv=data_dir+'/train.csv' # The training data only will be used for EDA, performance  
sample_submission=data_dir+'/sample_submission.csv'  
data_description='/data_description.txt'
```

```
train_df=pd.read_csv(train_csv)
```

```
train_df
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolA
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	...	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	...	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	...	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	...	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	...	
...	
1455	1456	60	RL	62.0	7917	Pave	NaN	Reg	Lvl	AllPub	...	
1456	1457	20	RL	85.0	13175	Pave	NaN	Reg	Lvl	AllPub	...	
1457	1458	70	RL	66.0	9042	Pave	NaN	Reg	Lvl	AllPub	...	
1458	1459	20	RL	68.0	9717	Pave	NaN	Reg	Lvl	AllPub	...	
1459	1460	20	RL	75.0	9937	Pave	NaN	Reg	Lvl	AllPub	...	

1460 rows × 81 columns

```
test_df=pd.read_csv(test_csv)  
test_df
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	Screen
--	----	------------	----------	-------------	---------	--------	-------	----------	-------------	-----------	-----	--------

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	Screen
0	1461	20	RH	80.0	11622	Pave	NaN	Reg	Lvl	AllPub	...	
1	1462	20	RL	81.0	14267	Pave	NaN	IR1	Lvl	AllPub	...	
2	1463	60	RL	74.0	13830	Pave	NaN	IR1	Lvl	AllPub	...	
3	1464	60	RL	78.0	9978	Pave	NaN	IR1	Lvl	AllPub	...	
4	1465	120	RL	43.0	5005	Pave	NaN	IR1	HLS	AllPub	...	
...	
1454	2915	160	RM	21.0	1936	Pave	NaN	Reg	Lvl	AllPub	...	
1455	2916	160	RM	21.0	1894	Pave	NaN	Reg	Lvl	AllPub	...	
1456	2917	20	RL	160.0	20000	Pave	NaN	Reg	Lvl	AllPub	...	
1457	2918	85	RL	62.0	10441	Pave	NaN	Reg	Lvl	AllPub	...	
1458	2919	60	RL	74.0	9627	Pave	NaN	Reg	Lvl	AllPub	...	

1459 rows × 80 columns

```
sample_submission_df=pd.read_csv(sample_submission)
sample_submission_df
```

	Id	SalePrice
0	1461	169277.052498
1	1462	187758.393989
2	1463	183583.683570
3	1464	179317.477511
4	1465	150730.079977
...
1454	2915	167081.220949
1455	2916	164788.778231
1456	2917	219222.423400
1457	2918	184924.279659
1458	2919	187741.866657

1459 rows × 2 columns

```
jovian.commit()
```

[jovian] Updating notebook "kavinm642/house-prices-prediction" on <https://jovian.com>
[jovian] Committed successfully! <https://jovian.com/kavinm642/house-prices-prediction>
'<https://jovian.com/kavinm642/house-prices-prediction>'

Exploratory Data Analysis

```
#checking the no.of rows and columns
n_rows=1460
```

```
n_cols=81
print('The dataset contains {} rows and {} columns.'.format(n_rows,n_cols))
```

The dataset contains 1460 rows and 81 columns.

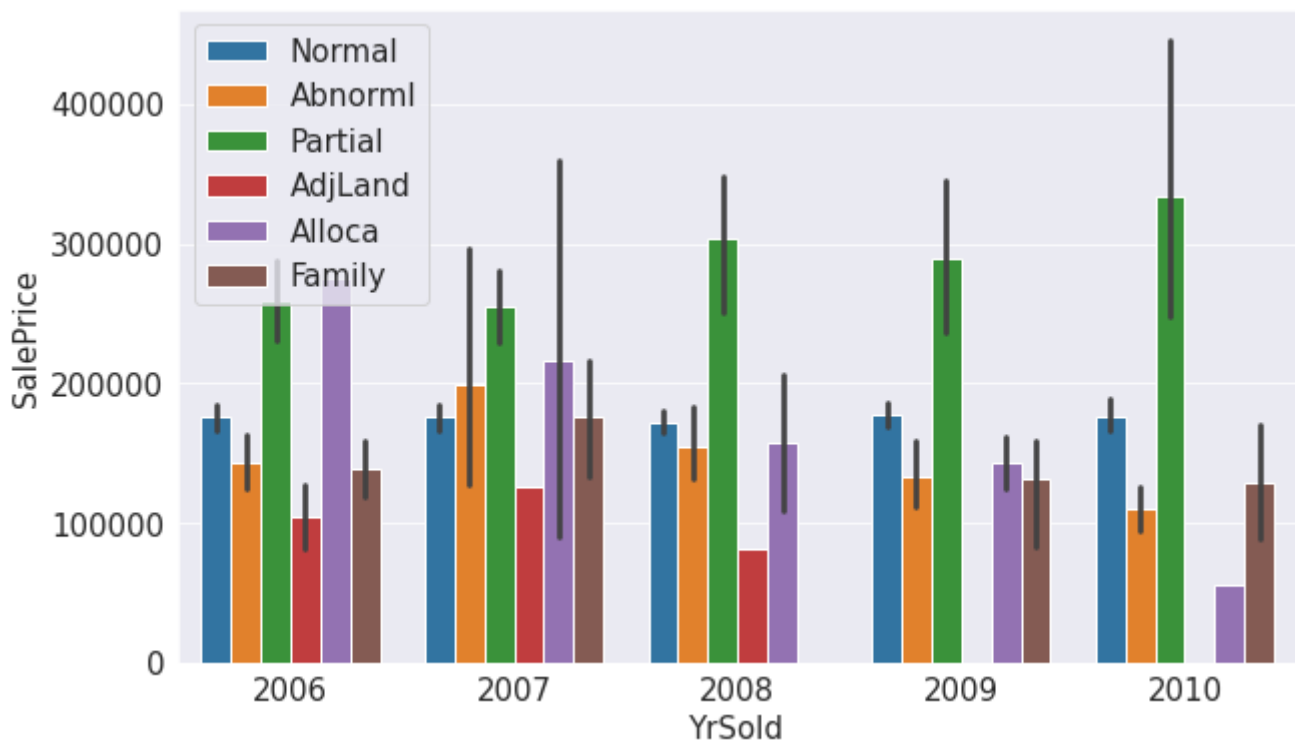
```
!pip install plotly
```

Requirement already satisfied: plotly in /opt/conda/lib/python3.9/site-packages (5.13.0)

Requirement already satisfied: tenacity>=6.2.0 in /opt/conda/lib/python3.9/site-packages (from plotly) (8.1.0)

```
# For visualizing the data
import plotly.express as px
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib
%matplotlib inline
sns.set_style('darkgrid')
matplotlib.rcParams['font.size']=15
matplotlib.rcParams['figure.figsize']=(10,6)
```

```
# year sold and sale price
sns.barplot(data=train_df,x='YrSold',y='SalePrice',hue='SaleCondition')
plt.legend(loc='upper left')
plt.show()
```



```
#comparing the amenities of Garage in a house according to its price
fig=px.scatter(train_df,x='GarageYrBlt',y='GrLivArea',color='GarageType',opacity=1.0,ho
```

```
fig.update_traces(marker_size=5)
fig.show()
```

```
train_df.BsmtFinSF1.corr(train_df.BsmtFinSF2)
```

```
-0.050117400047150915
```

```
train_df.LotFrontage.corr(train_df.LotArea)
```

```
0.4260950187718078
```

```
jovian.commit()
```

```
[jovian] Updating notebook "kavinm642/house-prices-prediction" on https://jovian.com
[jovian] Committed successfully! https://jovian.com/kavinm642/house-prices-prediction
'https://jovian.com/kavinm642/house-prices-prediction'
```

Preparing the dataset for training

```
train_df
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolA
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	...	

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolA
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	...	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	...	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	...	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	...	
...	
1455	1456	60	RL	62.0	7917	Pave	NaN	Reg	Lvl	AllPub	...	
1456	1457	20	RL	85.0	13175	Pave	NaN	Reg	Lvl	AllPub	...	
1457	1458	70	RL	66.0	9042	Pave	NaN	Reg	Lvl	AllPub	...	
1458	1459	20	RL	68.0	9717	Pave	NaN	Reg	Lvl	AllPub	...	
1459	1460	20	RL	75.0	9937	Pave	NaN	Reg	Lvl	AllPub	...	

1460 rows × 81 columns

`input_col` which contains the data that can be used as an input to train the model.

`target_col` which contains the data that can be used as a target to train the model.

```
input_col=['MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street', 'Alley', 'LotS
```

```
target_col='SalePrice'
```

```
print(list(input_col))
```

```
['MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street', 'Alley', 'LotShape',
'LandContour', 'Utilities', 'LotConfig', 'LandSlope', 'Neighborhood', 'Condition1',
'Condition2', 'BldgType', 'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt',
'YearRemodAdd', 'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond',
'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1', 'BsmtFinType2', 'BsmtFinSF2',
'BsmtUnfSF', 'TotalBsmtSF', 'Heating', 'HeatingQC', 'CentralAir', 'Electrical',
'1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath',
'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual', 'TotRmsAbvGrd',
'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType', 'GarageYrBlt', 'GarageFinish',
'GarageCars', 'GarageArea', 'GarageQual', 'GarageCond', 'PavedDrive', 'WoodDeckSF',
'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',
'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType', 'SaleCondition']
```

```
len(input_col)
```

79

```
print(target_col)
```

SalePrice

```
input_df=train_df[input_col].copy()
```

```
target_df=train_df[target_col].copy()
```

input_df

	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	LotConfig	...	\$
0	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	Inside	...	
1	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	FR2	...	
2	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	Inside	...	
3	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	Corner	...	
4	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	FR2	...	
...	
1455	60	RL	62.0	7917	Pave	NaN	Reg	Lvl	AllPub	Inside	...	
1456	20	RL	85.0	13175	Pave	NaN	Reg	Lvl	AllPub	Inside	...	
1457	70	RL	66.0	9042	Pave	NaN	Reg	Lvl	AllPub	Inside	...	
1458	20	RL	68.0	9717	Pave	NaN	Reg	Lvl	AllPub	Inside	...	
1459	20	RL	75.0	9937	Pave	NaN	Reg	Lvl	AllPub	Inside	...	

1460 rows × 79 columns

target_df

```
0      208500
1      181500
2      223500
3      140000
4      250000
...
1455    175000
1456    210000
1457    266500
1458    142125
1459    147500
```

Name: SalePrice, Length: 1460, dtype: int64

```
jovian.commit()
```

[jovian] Updating notebook "kavinm642/house-prices-prediction" on <https://jovian.com>
[jovian] Committed successfully! <https://jovian.com/kavinm642/house-prices-prediction>
'<https://jovian.com/kavinm642/house-prices-prediction>'

Identifying Numeric and Categorical data


```
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1460 entries, 0 to 1459
```

```
Data columns (total 81 columns):
```

#	Column	Non-Null Count	Dtype
0	Id	1460 non-null	int64
1	MSSubClass	1460 non-null	int64
2	MSZoning	1460 non-null	object
3	LotFrontage	1201 non-null	float64
4	LotArea	1460 non-null	int64
5	Street	1460 non-null	object
6	Alley	91 non-null	object
7	LotShape	1460 non-null	object
8	LandContour	1460 non-null	object
9	Utilities	1460 non-null	object
10	LotConfig	1460 non-null	object
11	LandSlope	1460 non-null	object
12	Neighborhood	1460 non-null	object
13	Condition1	1460 non-null	object
14	Condition2	1460 non-null	object
15	BldgType	1460 non-null	object
16	HouseStyle	1460 non-null	object
17	OverallQual	1460 non-null	int64
18	OverallCond	1460 non-null	int64
19	YearBuilt	1460 non-null	int64
20	YearRemodAdd	1460 non-null	int64
21	RoofStyle	1460 non-null	object
22	RoofMatl	1460 non-null	object
23	Exterior1st	1460 non-null	object
24	Exterior2nd	1460 non-null	object
25	MasVnrType	1452 non-null	object
26	MasVnrArea	1452 non-null	float64
27	ExterQual	1460 non-null	object
28	ExterCond	1460 non-null	object
29	Foundation	1460 non-null	object
30	BsmtQual	1423 non-null	object
31	BsmtCond	1423 non-null	object
32	BsmtExposure	1422 non-null	object
33	BsmtFinType1	1423 non-null	object
34	BsmtFinSF1	1460 non-null	int64
35	BsmtFinType2	1422 non-null	object

36	BsmtFinSF2	1460	non-null	int64
37	BsmtUnfSF	1460	non-null	int64
38	TotalBsmtSF	1460	non-null	int64
39	Heating	1460	non-null	object
40	HeatingQC	1460	non-null	object
41	CentralAir	1460	non-null	object
42	Electrical	1459	non-null	object
43	1stFlrSF	1460	non-null	int64
44	2ndFlrSF	1460	non-null	int64
45	LowQualFinSF	1460	non-null	int64
46	GrLivArea	1460	non-null	int64
47	BsmtFullBath	1460	non-null	int64
48	BsmtHalfBath	1460	non-null	int64
49	FullBath	1460	non-null	int64
50	HalfBath	1460	non-null	int64
51	BedroomAbvGr	1460	non-null	int64
52	KitchenAbvGr	1460	non-null	int64
53	KitchenQual	1460	non-null	object
54	TotRmsAbvGrd	1460	non-null	int64
55	Functional	1460	non-null	object
56	Fireplaces	1460	non-null	int64
57	FireplaceQu	770	non-null	object
58	GarageType	1379	non-null	object
59	GarageYrBltd	1379	non-null	float64
60	GarageFinish	1379	non-null	object
61	GarageCars	1460	non-null	int64
62	GarageArea	1460	non-null	int64
63	GarageQual	1379	non-null	object
64	GarageCond	1379	non-null	object
65	PavedDrive	1460	non-null	object
66	WoodDeckSF	1460	non-null	int64
67	OpenPorchSF	1460	non-null	int64
68	EnclosedPorch	1460	non-null	int64
69	3SsnPorch	1460	non-null	int64
70	ScreenPorch	1460	non-null	int64
71	PoolArea	1460	non-null	int64
72	PoolQC	7	non-null	object
73	Fence	281	non-null	object
74	MiscFeature	54	non-null	object
75	MiscVal	1460	non-null	int64
76	MoSold	1460	non-null	int64
77	YrSold	1460	non-null	int64
78	SaleType	1460	non-null	object

```
79 SaleCondition 1460 non-null object
80 SalePrice      1460 non-null int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB
```

```
numeric_col=input_df.select_dtypes(include=['int64','float64']).columns.tolist()
```

```
categoric_col=input_df.select_dtypes(include=['object']).columns.tolist()
```

```
print(list(numeric_col))
```

```
['MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual', 'OverallCond', 'YearBuilt',
'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF',
'1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath',
'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces',
'GarageYrBlt', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF',
'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal', 'MoSold', 'YrSold']
```

```
print(list(categoric_col))
```

```
['MSZoning', 'Street', 'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle',
'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'ExterQual',
'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1',
'BsmtFinType2', 'Heating', 'HeatingQC', 'CentralAir', 'Electrical', 'KitchenQual',
'Functional', 'FireplaceQu', 'GarageType', 'GarageFinish', 'GarageQual', 'GarageCond',
'PavedDrive', 'PoolQC', 'Fence', 'MiscFeature', 'SaleType', 'SaleCondition']
```

Impute Numerical Data

```
missing_count=input_df[numeric_col].isna().sum().sort_values(ascending=False)
missing_count[missing_count>0]
```

```
LotFrontage    259
GarageYrBlt     81
MasVnrArea       8
dtype: int64
```

Importing Simple Imputer

```
from sklearn.impute import SimpleImputer
```

```
imputer=SimpleImputer(strategy='mean')
```

```
imputer.fit(input_df[numeric_col])
```

```
SimpleImputer()
```

```
input_df[numeric_col]=imputer.transform(input_df[numeric_col])
```

```
missing_count=input_df[numeric_col].isna().sum().sort_values(ascending=False)
missing_count[missing_count>0]
```

```
Series([], dtype: int64)
```

```
jovian.commit()
```

```
[jovian] Updating notebook "kavinm642/house-prices-prediction" on https://jovian.com
[jovian] Committed successfully! https://jovian.com/kavinm642/house-prices-prediction
'https://jovian.com/kavinm642/house-prices-prediction'
```

Scaling Numerical Values

```
input_df[numeric_col].describe().loc[['min', 'max']]
```

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtFinS
min	20.0	21.0	1300.0	1.0	1.0	1872.0	1950.0	0.0	0.0
max	190.0	313.0	215245.0	10.0	9.0	2010.0	2010.0	1600.0	564.0

2 rows × 36 columns

Importing MinMaxScaler

```
scaler=MinMaxScaler()
```

```
scaler.fit(input_df[numeric_col])
```

```
MinMaxScaler()
```

```
input_df[numeric_col]=scaler.transform(input_df[numeric_col])
```

```
input_df[numeric_col].describe().loc[['min', 'max']]
```

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtFinSF1
min	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
max	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0

2 rows × 36 columns

Here the values turns into boolean expressions after scaling.

Encoding Categorical Values

```
input_df[categoric_col].nunique().sort_values(ascending=False)
```

Neighborhood	25
Exterior2nd	16
Exterior1st	15
SaleType	9
Condition1	9
Condition2	8
HouseStyle	8
RoofMatl	8
Functional	7
BsmtFinType2	6
Heating	6
RoofStyle	6
SaleCondition	6
BsmtFinType1	6
GarageType	6
Foundation	6
Electrical	5
FireplaceQu	5
HeatingQC	5
GarageQual	5
GarageCond	5
MSZoning	5
LotConfig	5
ExterCond	5
BldgType	5
BsmtExposure	4
MiscFeature	4
Fence	4
LotShape	4
LandContour	4
BsmtCond	4
KitchenQual	4
MasVnrType	4
ExterQual	4
BsmtQual	4
LandSlope	3
GarageFinish	3
PavedDrive	3
PoolQC	3
Utilities	2
CentralAir	2
Street	2
Alley	2

dtype: int64

Importing OneHotEncoder

```
from sklearn.preprocessing import OneHotEncoder
```

```
encoder=OneHotEncoder(sparse=False,handle_unknown='ignore')
```

```
encoder.fit(input_df[categorical_col])
```

```
OneHotEncoder(handle_unknown='ignore', sparse=False)
```

```
encoded_col=list(encoder.get_feature_names_out(categorical_col))  
len(encoded_col)
```

268

```
input_df[encoded_col]=encoder.transform(input_df[categorical_col])
```

/opt/conda/lib/python3.9/site-packages/pandas/core/frame.py:3678: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling `frame.insert` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use `newframe = frame.copy()`

input_df

	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	LotConfig	...
0	0.235294	RL	0.150685	0.033420	Pave	NaN	Reg	Lvl	AllPub	Inside	...
1	0.000000	RL	0.202055	0.038795	Pave	NaN	Reg	Lvl	AllPub	FR2	...
2	0.235294	RL	0.160959	0.046507	Pave	NaN	IR1	Lvl	AllPub	Inside	...
3	0.294118	RL	0.133562	0.038561	Pave	NaN	IR1	Lvl	AllPub	Corner	...
4	0.235294	RL	0.215753	0.060576	Pave	NaN	IR1	Lvl	AllPub	FR2	...
...
1455	0.235294	RL	0.140411	0.030929	Pave	NaN	Reg	Lvl	AllPub	Inside	...
1456	0.000000	RL	0.219178	0.055505	Pave	NaN	Reg	Lvl	AllPub	Inside	...
1457	0.294118	RL	0.154110	0.036187	Pave	NaN	Reg	Lvl	AllPub	Inside	...
1458	0.000000	RL	0.160959	0.039342	Pave	NaN	Reg	Lvl	AllPub	Inside	...
1459	0.000000	RL	0.184932	0.040370	Pave	NaN	Reg	Lvl	AllPub	Inside	...

1460 rows × 347 columns

```
jovian.commit()
```

[jovian] Updating notebook "kavinm642/house-prices-prediction" on <https://jovian.com>
[jovian] Committed successfully! <https://jovian.com/kavinm642/house-prices-prediction>
'<https://jovian.com/kavinm642/house-prices-prediction>'

Training and Validation Set

```
x_train,x_test,y_train,y_test=train_test_split(input_df[numeric_col+encoded_col],target
```

x_train

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtFin
1023	0.588235	0.075342	0.008797	0.666667	0.500	0.963768	0.933333	0.008750	0.002
810	0.000000	0.195205	0.041319	0.555556	0.625	0.739130	0.816667	0.061875	0.117
1384	0.176471	0.133562	0.036271	0.555556	0.500	0.485507	0.000000	0.000000	0.036
626	0.000000	0.167979	0.051611	0.444444	0.500	0.637681	0.466667	0.000000	0.000
813	0.000000	0.184932	0.039496	0.555556	0.625	0.623188	0.133333	0.151875	0.107
...
1095	0.000000	0.195205	0.037472	0.555556	0.500	0.971014	0.933333	0.000000	0.004
1130	0.176471	0.150685	0.030400	0.333333	0.250	0.405797	0.000000	0.000000	0.110
1294	0.000000	0.133562	0.032120	0.444444	0.750	0.601449	0.666667	0.000000	0.029
860	0.176471	0.116438	0.029643	0.666667	0.875	0.333333	0.800000	0.000000	0.000
1126	0.588235	0.109589	0.011143	0.666667	0.500	0.978261	0.950000	0.081250	0.000

1095 rows × 304 columns

x_test

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtFin
892	0.000000	0.167808	0.033252	0.555556	0.875	0.659420	0.883333	0.000000	0.117
1105	0.235294	0.263699	0.051209	0.777778	0.500	0.884058	0.750000	0.226250	0.182
413	0.058824	0.119863	0.035804	0.444444	0.625	0.398551	0.000000	0.000000	0.000
522	0.176471	0.099315	0.017294	0.555556	0.750	0.543478	0.000000	0.000000	0.070
1036	0.000000	0.232877	0.054210	0.888889	0.500	0.978261	0.966667	0.043750	0.181
...
988	0.235294	0.167979	0.050228	0.555556	0.625	0.753623	0.433333	0.186250	0.027
243	0.823529	0.184932	0.044226	0.555556	0.625	0.782609	0.500000	0.000000	0.000
1342	0.235294	0.167979	0.037743	0.777778	0.500	0.942029	0.866667	0.093125	0.000
1057	0.235294	0.167979	0.133955	0.666667	0.625	0.884058	0.733333	0.000000	0.105
1418	0.000000	0.171233	0.036944	0.444444	0.500	0.659420	0.216667	0.000000	0.004

365 rows × 304 columns

y_train

1023	191000
810	181000
1384	105000

```
626      139900
813      157900
...
1095     176432
1130     135000
1294     115000
860      189950
1126     174000
```

Name: SalePrice, Length: 1095, dtype: int64

```
y_test
```

```
892      154500
1105     325000
413      115000
522      159000
1036     315500
...
988      195000
243      120000
1342     228500
1057     248000
1418     124000
```

Name: SalePrice, Length: 365, dtype: int64

Random Forest Classifier

```
random=RandomForestClassifier(n_jobs=-1, random_state=42, n_estimators=10)
```

```
random.fit(x_train, y_train)
```

```
RandomForestClassifier(n_estimators=10, n_jobs=-1, random_state=42)
```

```
y_pred=random.predict(x_test)
```

Making predictions and evaluating the model

The model's performance will be evaluated using **RMSE (Root Mean Squared Error)** loss function.

It is one of the main performance indicator for a regression model. It measures the average difference between values predicted by a model and the actual values. It provides an estimation of how well the model is able to predict the target value (accuracy).

```
from sklearn.metrics import mean_squared_error
```

```
train_pred=random.predict(x_train)
```

```
train_rmse=random.score(x_train, y_train)
```



```
train_rmse #train attributes(train_inputs and train targets)
```

```
0.9963470319634703
```

```
print('The RMSE loss for the training set is $ {}'.format(train_rmse))
```

The RMSE loss for the training set is \$ 0.9963470319634703.

```
val_pred=y_test
```

```
val_rmse=random.score(x_test,y_test)
```

```
val_rmse #target [validation] attributes (val_inputs and val_targets)
```

```
0.005479452054794521
```

```
print('The RMSE loss for the validation set is $ {}'.format(val_rmse))
```

The RMSE loss for the validation set is \$ 0.005479452054794521.

```
jovian.commit()
```

[jovian] Updating notebook "kavinm642/house-prices-prediction" on <https://jovian.com>
[jovian] Committed successfully! <https://jovian.com/kavinm642/house-prices-prediction>
'<https://jovian.com/kavinm642/house-prices-prediction>'

Making Predictions

```
def predict_input(single_input):  
    input_df = pd.DataFrame([single_input])  
    input_df[numeric_col] = imputer.transform(input_df[numeric_col])  
    input_df[numeric_col] = scaler.transform(input_df[numeric_col])  
    input_df[encoded_col] = encoder.transform(input_df[categorical_col].values)  
    X_input = input_df[numeric_col + encoded_col]  
    return random.predict(X_input)[0]
```

```
sample_input = { 'MSSubClass': 20, 'MSZoning': 'RL', 'LotFrontage': 77.0, 'LotArea': 93  
'Street': 'Pave', 'Alley': None, 'LotShape': 'IR1', 'LandContour': 'Lvl', 'Utilities':  
'LotConfig': 'Inside', 'LandSlope': 'Gtl', 'Neighborhood': 'Names', 'Condition1': 'Nor  
'BldgType': '1Fam', 'HouseStyle': '1Story', 'OverallQual': 4, 'OverallCond': 5, 'YearB  
'YearRemodAdd': 1959, 'RoofStyle': 'Gable', 'RoofMatl': 'CompShg', 'Exterior1st': 'Ply  
'Exterior2nd': 'Plywood', 'MasVnrType': 'None', 'MasVnrArea': 0.0, 'ExterQual': 'TA', 'Ex  
'Foundation': 'CBlock', 'BsmtQual': 'TA', 'BsmtCond': 'TA', 'BsmtExposure': 'No', 'BsmtFin  
'BsmtFinSF1': 569, 'BsmtFinType2': 'Unf', 'BsmtFinSF2': 0, 'BsmtUnfSF': 381,  
'TotalBsmtSF': 950, 'Heating': 'GasA', 'HeatingQC': 'Fa', 'CentralAir': 'Y', 'Electrical':  
'2ndFlrSF': 0, 'LowQualFinSF': 0, 'GrLivArea': 1225, 'BsmtFullBath': 1, 'BsmtHalfBath'  
'HalfBath': 1, 'BedroomAbvGr': 3, 'KitchenAbvGr': 1, 'KitchenQual': 'TA', 'TotRmsAbvGrd'
```

```
'Fireplaces': 0, 'FireplaceQu': np.nan, 'GarageType': np.nan, 'GarageYrBlt': np.nan, 'GarageArea': 0, 'GarageQual': np.nan, 'GarageCond': np.nan, 'PavedDrive': 'Y', 'WoodDeck': 0, 'EnclosedPorch': 0, '3SsnPorch': 0, 'ScreenPorch': 0, 'PoolArea': 0, 'PoolQC': np.nan, 'MiscVal': 400, 'MoSold': 1, 'YrSold': 2010, 'SaleType': 'WD', 'SaleCondition': 'Normal'
```

```
predicted_price=predict_input(sample_input)
```

/opt/conda/lib/python3.9/site-packages/sklearn/base.py:445: UserWarning:

X does not have valid feature names, but OneHotEncoder was fitted with feature names

/opt/conda/lib/python3.9/site-packages/pandas/core/frame.py:3678: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling `frame.insert` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use `newframe = frame.copy()`

```
print('The predicted sale price of the house is ${}.'.format(predicted_price))
```

The predicted sale price of the house is \$52500.

Conclusion

Thus, the final price of a house is **\$52500**. However the cost will vary according to the changes committing in `sample_input`

```
jovian.commit()
```

```
[jovian] Updating notebook "kavinm642/house-prices-prediction" on https://jovian.com  
[jovian] Committed successfully! https://jovian.com/kavinm642/house-prices-prediction  
'https://jovian.com/kavinm642/house-prices-prediction'
```